

Economic Integrity in Production LLM Pipelines: Measuring Attempt Amplification, Discarded Reasoning, and Retry Cost

Sean Halverson
VuduVations · vuduvations.io
hello@vuduvations.io

August 2026 (preprint v0.2.4)

Patent pending. One or more U.S. provisional patent applications have been filed covering aspects of the systems and methods described herein.

Abstract

Model pricing pages quote cost per token. Production pipelines pay per *attempt*. When a pipeline retries empty replies, when its validation layers re-ask failed roles, and when reasoning models bill thinking tokens for outputs they never emit, the invoice decouples from the scoreboard. In a companion reliability study we reported this decoupling as an instrumentation gap: attempts were not instrumented, so attempt denominators were unavailable and starvation could only be counted from run logs. This paper closes that gap. We add a per-attempt economic ledger at the HTTP boundary of a production analysis pipeline — one record per physical model call, carrying token counts including reasoning and prompt-cache splits, budgets, lever state, layer, outcome, and latency — and re-run a four-cell budget/retry ablation ($n=10$ repetitions per cell, one provider, one night) with the ledger armed. We define two metrics over the ledger: the *Inference Amplification Factor* (IAF), physical attempts per logical job, and the *Economic Contamination Index* (ECI), total spend over spend on attempts whose output was actually used. All four cells were scoreboard-equivalent — answer-key means of 46.8–47.4 of 48 checks and 9–10 of 10 execution-clean or recovered-clean repetitions; underneath, the attempt machinery differed by up to a factor of 1.7. Proactive budget headroom *measures* as prevention: one starved first attempt in 127 agent-layer jobs, and — with the retry disabled, so no recovery could hide it — no amplification at either recovery ring (IAF = 1.000 by direct enumeration). Retry-based recovery reaches an equivalent scoreboard with IAF up to 1.70 and ECI up to 1.96 in the adversarial validation layer: in the worst cell, 48.9% of that layer’s spend bought attempts whose output was discarded — and 228,889 of 228,890 discarded output tokens in the study were reasoning tokens, the spend category no scoreboard sees. Rescue effectiveness was also non-stationary: a same-budget re-ask recovered roughly 19 of 20 starved jobs one night and 5 of 16 two nights later under an identical configuration, so retry-based cost forecasts inherit a volatility that prevention-based forecasts do not. Across the study, 9.6% of ledger-computed spend bought discarded attempts. A retry repairs the answer. It does not erase the first invoice.

1 Introduction

Benchmark papers report accuracy per task. Pricing pages report dollars per million tokens. Between the two sits an object neither one measures: the *attempt*. A production pipeline that

values availability will retry an empty reply, re-ask a failed validation role, and salvage what it can; each of those recovery actions is a fresh, billed model call whose predecessor’s cost is already sunk. When the model family is a reasoning model, the sunk cost is worse than it looks: an attempt that exhausts its completion budget on chain-of-thought and emits zero visible output still bills its reasoning tokens at the output rate [2, 3]. The scoreboard records one answer. The invoice records every attempt. Nothing in the standard reporting stack reconciles the two.

Our companion reliability study [1] documented the epistemic half of this problem: a fallback-enabled pipeline can produce near-perfect benchmark scores while the probabilistic component under test did not produce them, and run-level *degradation accounting* is required to tell the difference. That paper counted starvation *events* from run logs and was explicit about its boundary: attempts were not instrumented, so it declined attempt denominators and reported counts, not rates. The present paper is the economic half. We instrument the attempt itself and ask what the recovery machinery *costs* — not in the aggregate-bill sense, but attributed per attempt, per layer, and per outcome.

The contributions are:

1. **An attempt-level economic ledger** (Section 4): one JSONL record per physical HTTP model call, written inside the pipeline’s portable client below all metering wrappers, carrying input, output, and reasoning token counts, prompt-cache hit/miss splits, requested and sent budgets, intervention-lever state, pipeline layer, attempt number, outcome, and latency. The hook is armed by an environment variable and is a strict no-op otherwise; the pipeline’s legacy per-run token counters are deliberately unchanged, so the ledger-versus-legacy delta is itself a measurement of what conventional metering misses.
2. **Two metrics with measured denominators** (Section 3): the Inference Amplification Factor (IAF), physical attempts per logical job, computed per pipeline layer and never pooled across layers; and the Economic Contamination Index (ECI), total spend over spend on used attempts, where *used* is defined exactly (the final attempt of a logical job, when non-empty).
3. **A measured four-cell ablation** (Sections 5–6): the companion paper’s budget/retry ablation re-run with the ledger armed ($n=10$ per cell, DeepSeek v4-pro direct, 2026-08-10/11). Every cell was scoreboard-equivalent (answer-key means 46.8–47.4 of 48; 9–10 of 10 execution-clean or recovered-clean repetitions); the cells’ IAF and ECI separated cleanly, and in the direction the companion paper’s event counts predicted.
4. **Three economic findings**: prevention is measurable by enumeration — a counted starvation residual with no hidden recovery amplification — rather than inferable from absence of log lines (Section 7); the adversarial validation layer, whose role calls run on small completion budgets, was the study’s dominant amplifier, with up to half its spend discarded (Section 8); and rescue rates moved by a factor of about three across two nights under identical configuration, which bounds how retry-heavy architectures can be cost-forecast (Section 9).

We state the scope plainly, as in the companion study: this is a measurement-method case study on one production pipeline, one provider, one night per cell, $n=10$ per cell. The numbers are counts and dollars with denominators we measured, not rates we expect to generalize. What we believe generalizes is the instrument and the two metrics.

2 The Hidden Denominator Problem

Three mechanisms make the attempt invisible to conventional accounting.

Retries rebind the response. A recovery path that retries an empty reply typically overwrites the failed response object with the successful one; whatever token metering reads the response afterward meters only the survivor. Our own pipeline had exactly this seam: the portable client’s empty-reply retry rebound the response variable, so the first attempt’s tokens — a starved attempt bills its entire reasoning budget — were counted by nothing. The pipeline’s per-run counters, and every figure derived from them, silently reported the cost of the attempts that *worked*.

Reasoning tokens are invisible output. On reasoning-mode models, chain-of-thought is billed at the output rate but is not part of the visible completion [2, 3]. A starved attempt — one that spends its whole completion budget thinking and emits nothing — therefore produces a bill with no corresponding artifact anywhere in the application. In this study, 65.1% of all billed output tokens were reasoning tokens, and 228,889 of 228,890 discarded output tokens were reasoning tokens: the attempts the pipeline threw away consisted almost entirely of thinking no one ever saw.

Layered recovery multiplies callers. Production pipelines recover at more than one level. In ours, the portable client retries an empty reply once; above it, the adversarial validation layer’s orchestrator can re-ask a failed role as a fresh call. Each ring of recovery bills separately, and an aggregate invoice cannot attribute spend to rings. Only an attempt-level record with a logical-job identity can.

The result is a denominator problem. “Cost per task” has a measured numerator (the bill) and an assumed denominator (one attempt per task). The companion study showed the assumption failing at rates up to 96 starvation events per ten repetitions [1] — but could not price it, because the failed attempts were not instrumented. The remainder of this paper is the pricing.

3 The Economic Integrity Framework

Three layers of integrity. A production result can be interrogated at three depths. *Capability integrity*: did the answer satisfy the task? This is what the benchmark’s 48-check answer key grades. *Execution integrity*: did the intended model path produce it? This is the companion paper’s degradation accounting [1] — fallbacks, salvages, and retries that leave the answer intact while changing what produced it. *Economic integrity*: how many paid physical attempts stood behind the accepted output? Each layer can pass while the next fails silently: a correct answer can come from a fallback, and a fallback-free run can have paid for twice the attempts its architecture diagram shows. This paper supplies the third layer’s instrument and metrics; the four cells of Section 6 are precisely a set of configurations that tie on the first layer, mostly tie on the second, and separate on the third.

We define the metrics over an attempt ledger: a set of records, one per physical model call, each carrying a logical-job identity, an attempt number within that job, token counts, and an outcome.

Logical job. One unit of work the pipeline intended: one agent completion, one validation-role completion. A logical job is identified by its calling site and a per-run sequence number, so two calls from the same site are distinct jobs, and a retry of one call is the same job.

Physical attempt. One HTTP round trip that was billed (or errored). A logical job has one or more physical attempts.

Inference Amplification Factor. For a set of logical jobs J with physical attempts A ,

$$\text{IAF} = \frac{|A|}{|J|}.$$

IAF = 1 means every job succeeded on its first attempt; IAF = 1.5 means half again as many calls were made as the architecture diagram shows. We compute IAF *per pipeline layer* (primary agents versus the adversarial validation layer) and never pool layers: the layers run different budgets, different prompt shapes, and different recovery policies, and a pooled figure would average away exactly the structure the metric exists to expose.

Used attempt. The final attempt of a logical job, when its content is non-empty. Every other attempt — an empty attempt that was retried, or a final attempt that was itself empty — was paid for and discarded. This definition is deliberately conservative in one known way: the pipeline has salvage paths that can extract partial value from a malformed reply, and the ledger classifies such a reply by emptiness, not by downstream salvage. ECI as computed here is therefore a mild upper bound on contamination.

Economic Contamination Index. For a layer with total attempt spend C_{total} and spend on used attempts C_{used} ,

$$\text{ECI} = \frac{C_{\text{total}}}{C_{\text{used}}}.$$

ECI = 1 means every dollar bought output that was used; ECI = 2 means half the spend bought discarded attempts. ECI is the economic analogue of the companion paper’s fallback contamination: there, the scoreboard was contaminated by answers the model under test did not produce; here, the invoice is contaminated by attempts the application never used.

Both metrics have *measured* denominators: every physical attempt in this study is a ledger row, including attempts that errored at the transport layer. The companion paper’s caveat — “attempts were not instrumented” — does not apply to these runs.

4 Instrumentation

Placement. The ledger hook lives inside the pipeline’s portable LLM client, immediately around each transport-level completion call — below the benchmark’s metering wrapper and below the client’s own retry logic. Placement matters twice over. First, it sits under the response-rebinding seam described in Section 2, so the starved first attempt that legacy metering loses is a first-class row. Second, the adversarial validation layer’s role calls route through the same injected client in this configuration, so chamber traffic is captured by the same hook with no second instrument.

Record schema. Each row carries: timestamp; experiment cell and repetition; logical job id (calling agent name plus per-run sequence); layer (**agent** or **debate**, derived from the calling site, never pooled downstream); attempt number; requested and sent completion budgets; the state of the three intervention levers; input, output, and *reasoning* token counts from the provider’s usage object; prompt-cache hit and miss token splits; finish reason; serving identity where exposed; a content-emptiness flag; outcome (**ok**, **empty**, or **error**); and wall-clock latency. Rows are appended to a JSONL file named by an environment variable; when the variable is unset

the hook is a no-op, and the instrumented client shipped to production with the hook dormant. Ledger writes never fail a run: an unwritable ledger degrades to a warning.

The nested retry topology. The pipeline recovers at two rings, and the ledger records them differently, on purpose. The portable client’s own empty-reply retry is the *same logical job*, attempt number 2. The validation orchestrator’s re-ask of a failed role is a *new logical job*: it arrives at the client as a fresh call from the same calling site and takes the next sequence number. This is not a limitation but a feature of the identity design: ring membership is recoverable from the ledger (attempt numbers expose the inner ring; job-count inflation at a fixed calling-site population exposes the outer ring), and Section 8 uses exactly this signature — 193 debate-layer logical jobs in the cell whose inner retries mostly failed, against 110–119 in the others — to observe the outer ring engaging when the inner ring fails.

Legacy counters unchanged. The client’s pre-existing token counters, which meter only the returned attempt, were deliberately left untouched, for two reasons. They preserve comparability with the companion paper’s reported figures. And the delta between ledger and legacy totals is itself the headline measurement: it is precisely the paid-but-discarded spend that conventional metering cannot see.

Pricing basis. Dollar figures computed from the ledger use the provider’s undiscounted direct list rates (\$0.435 input, \$0.87 output per million tokens [3]), with cache-hit input tokens priced at the full input rate. Ledger-computed dollars are therefore upper bounds; hit/miss splits are recorded per attempt for re-pricing. Section 10 reconciles this basis against the account-billed total.

5 Experimental Design

The workload is the companion study’s fixed benchmark: four analysis protocols over a frozen fictional corpus, graded by a 48-check answer key, roughly 12–13 agent-layer completions and a multi-role adversarial validation pass per repetition [1]. The model is DeepSeek v4-pro on the direct first-party channel, the same model and channel as the companion paper’s post-fix runs. The manipulated factors are the same two mechanisms as the companion ablation, crossed into four cells:

Cell	Proactive headroom	Empty-reply retry
<i>reactive</i>	off	on, retry adds headroom
<i>proactive</i>	on	off
<i>pure retry</i>	off	on, retry at unchanged budget
<i>bundle</i>	on	on, retry adds headroom

Table 1: The four cells. *Proactive headroom* raises completion budgets before the first attempt; the *reactive* and *bundle* retries add the same headroom only on the second attempt; the *pure retry* cell re-asks at the unchanged budget, isolating re-sampling from budget relief. Cell configuration is applied by process environment at launch; the run-of-record script clears all levers between cells.

Each cell ran $n=10$ repetitions (five launches of two repetitions each, matching the companion run-of-record shape) on 2026-08-10/11, with the ledger armed and a balance-floor guard between cells. Unlike the companion ablation, whose cells were launched by hand, the entire run is a

committed script; the raw ledgers (one JSONL per cell), the twenty per-launch result files, and the compiled report are all committed alongside it (Section 12).

One asymmetry with the companion paper must be stated up front: these are *new* runs on a *different night*, not a re-analysis of the companion cells. The companion ablation’s agent-layer starvation ran 96–102 events per no-headroom cell on 2026-08-09; tonight the same configurations produced 12–16. Cross-night variance of this size is itself one of the paper’s findings (Section 9), and it is why every figure below carries its own denominators and no figure is offered as a rate.

6 Results

Table 2 is the study’s central result. All four cells were scoreboard-equivalent, in the two senses Table 3 makes precise (we use the term descriptively, not as a formal equivalence test: exact answer-key totals differ by cell and $n=10$ supports no stronger claim than that neither measurement orders the cells). On capability, answer-key totals ran 468–474 of 480 checks (cell means 46.8–47.4 of 48), and the differences do not order the cells — the only cell with 10 of 10 execution-clean-or-recovered repetitions has the *lowest* answer-key mean. On execution, every cell produced 9–10 of 10 repetitions that were execution-clean or recovered-clean. The attempt machinery underneath was not equivalent.

Cell	Layer	Jobs	Att.	IAF	Starved	Rec.	Unrec.	Total \$	Used \$	Disc. \$	ECI
reactive	agent	126	137	1.087	12	11	1	0.643	0.628	0.015	1.023
	debate	119	202	1.698	83	82	1	0.289	0.184	0.106	1.576
proactive	agent	127	127	1.000	1	0	1	0.680	0.680	0.000	1.000
	debate	119	119	1.000	1	0	1	0.190	0.190	0.001	1.003
pure retry	agent	128	142	1.109	16	5	11	0.652	0.623	0.029	1.046
	debate	193	281	1.456	88	19	69	0.380	0.194	0.186	1.957
bundle	agent	124	126	1.016	2	2	0	0.689	0.670	0.020	1.029
	debate	110	110	1.000	0	0	0	0.171	0.171	0.000	1.000

Table 2: Measured attempt economics per cell and layer ($n=10$ repetitions per cell, DeepSeek v4-pro direct, 2026-08-10/11). *Jobs* are logical jobs, *Att.* physical attempts, *Starved* first attempts returning empty, *Rec./Unrec.* starved jobs whose final attempt was non-empty/empty. Dollars are ledger-computed at undiscounted list rates with cache-hit input at the full input rate (upper bounds; Section 10). Layers are never pooled. Agent-layer job counts exceed the ten graded agents because summarizer and validator completions share the layer; the outcome tiers in Table 3 still grade on the ten-agent basis of the companion paper. The proactive agent layer’s single starved job was a zero-token transport error, hence discarded spend of \$0.000.

Four observations organize the rest of the paper.

(1) Scoreboard equivalence hid a $1.7\times$ attempt spread. The reactive cell’s debate layer made 202 physical calls to complete 119 logical jobs ($\text{IAF} = 1.698$); the proactive cell’s made 119 for 119. Both cells were scoreboard-equivalent (answer-key means 47.0 and 47.4; nine of ten repetitions clean or recovered). Nothing on either scoreboard, and nothing in legacy returned-attempt metering, distinguishes them.

Cell	Answer key (mean)	1st-att. clean	Recovered clean	Degraded	Agent \$/rep
reactive	470/480 (47.0)	0	9	1	0.064–0.067
proactive	474/480 (47.4)	9	0	1	0.065–0.068
pure retry	472/480 (47.2)	0	9	1	0.055–0.066
bundle	468/480 (46.8)	8	2	0	0.067–0.078

Table 3: Per-repetition scoreboards: answer-key totals (48-check benchmark, ten repetitions per cell) beside execution tiers (ten-agent grading basis joined to per-repetition ledger cost). The two are distinct measurements and neither orders the cells: the bundle is the only cell with zero degraded repetitions and also has the lowest answer-key mean. Mean agent-layer cost per repetition is nearly flat across cells and tiers at tonight’s event rates: the economic difference between the cells lives in Table 2’s debate rows and in attempt-count variance, not in the per-repetition agent bill.




Figure 1: The measured attempt economics of Table 2 ($n=10$ repetitions per cell, DeepSeek v4-pro direct, 2026-08-10/11; every value is read programmatically from the compiled report, none transcribed). Panels 1 and 2 show IAF and ECI per layer, never pooled; the dotted line at 1.0 is the no-amplification floor that both proactive rows sit on. Panel 3 splits the adversarial validation layer’s ledger-computed spend into used and discarded dollars: the two no-headroom cells discarded 37% and 49% of the layer’s bill to reach a scoreboard equivalent to the cells that discarded effectively none.

(2) Discarded spend concentrates where budgets are tight. The adversarial validation layer runs its roles on small completion budgets; without headroom it starved 83–88 times per cell and its ECI reached 1.58–1.96. The agent layer, on generous budgets tonight, held ECI at 1.02–1.05 everywhere. Across the whole study, \$0.355 of \$3.694 ledger-computed spend — 9.6% — bought attempts whose output was discarded, and 82% of that discarded spend was in the debate layer.

(3) Nearly every discarded output token was a reasoning token. Summed over all cells and layers, discarded attempts billed 228,890 output tokens, of which 228,889 were reasoning tokens. A starved attempt spends its entire completion budget thinking and emits nothing; nearly all discarded output tokens were invisible reasoning tokens. Overall, 65.1% of the study’s billed output tokens were reasoning tokens, so the spend category most exposed to discard is also the largest one.

(4) The ledger measured the outer recovery ring engaging. The pure-retry cell’s debate layer shows 193 logical jobs against 110–119 in every other cell. Its inner retries mostly

failed (19 of 88 starved jobs recovered at the unchanged budget), and the failures surfaced to the validation orchestrator, whose re-asks arrive as fresh logical jobs (Section 4). The job-count inflation is the outer ring’s billing signature: recovery that fails at one ring does not end the spending, it escalates it to the next ring’s budget.

7 Recovery Is Not Prevention, Priced

The companion paper’s ablation attributed *prevention* of first-attempt starvation to proactive headroom and characterized retry as insurance [1]. The ledger now prices that distinction, and closes a measurement blind spot in it.

Prevention measures as a counted residual with no hidden amplification. One caveat first: in the proactive cell the client retry is ablated, so inner-ring IAF cannot exceed 1 by construction — $\text{IAF} = 1.000$ is not, by itself, the evidence. The evidence is three measured facts the ledger adds. First, the starvation residual is enumerated, not inferred: 127 agent-layer jobs, 127 attempts, one empty reply (a zero-token transport error, a class no completion budget can prevent); 119 debate-layer jobs, 119 attempts, one starved. The companion paper’s “zero starvation events” claim had rested on the absence of a retry log line — a line that *cannot fire* in that cell, because the retry is the thing ablated; absence of evidence from a disabled instrument is not evidence of absence, and the ledger replaces that inference with enumeration. Second, the outer recovery ring did not engage: the validation orchestrator’s re-asks register as extra logical jobs (Section 4), and the debate-layer job count sits at the baseline 119, not the 193 of the cell where inner recovery failed. Third, the two residual empties went unrecovered — headroom without retry has no insurance, which is the bundle’s argument. The claim survives direct measurement — as 1-in-127, not zero, which is how such claims should be stated.

Recovery reaches an equivalent scoreboard and pays for both attempts. The reactive cell’s scoreboard was equivalent to proactive’s (nine of ten repetitions clean or recovered; answer-key means 47.0 versus 47.4). To get there its debate layer paid for 202 attempts against 119 jobs, discarding \$0.106 of \$0.289 — 36.5% of that layer’s bill — while the proactive cell’s debate layer discarded \$0.0006. On tonight’s calm agent layer the per-repetition bills were nearly identical across cells (Table 3); the economic difference between recovery and prevention lived in the tight-budget layer and in latency and variance, not in the agent-layer bill. On the companion paper’s night, when the agent layer starved 96–102 times per cell, the same structural difference would have priced an order of magnitude higher. That contrast is the forecasting point of Section 9: prevention’s cost is stable by construction (one attempt per job); recovery’s cost is a random variable.

Retry still earns its place. The bundle cell — headroom plus retry — was the only cell with zero degraded repetitions, and its two first-attempt agent starvations were both recovered by the retry at a discarded cost of \$0.020. As in the companion study, the argument is not against retry; it is against retry *as the mechanism* rather than as insurance behind prevention. The ledger’s version: insurance premiums (bundle, $\text{ECI} \leq 1.03$) are cheap; using insurance as the primary mechanism (pure retry, debate $\text{ECI} = 1.96$) doubles the effective price of the layer it protects.

8 The Adversarial-Layer Amplifier

The companion paper’s starvation story was agent-layer: 96–102 events per no-headroom cell on its night. Tonight the agent layer was calm (12–16 events per no-headroom cell, same configuration, same provider) and the amplification story moved down a layer. The adversarial validation chamber runs each debate role on a small completion budget — an economical choice for short verdicts, and precisely the budget most exposed to reasoning-mode starvation. Without headroom it starved on more than two-thirds of its jobs (83 of 119, 88 of 193), and because a starved role attempt bills its whole budget as reasoning tokens, the layer’s ECI ran 1.58 (reactive) and 1.96 (pure retry): in the pure-retry cell, 48.9% of the validation layer’s spend bought discarded attempts.

Two structural facts make small-budget validation layers the natural amplifier. First, the starvation probability is a function of budget versus the model’s reasoning appetite, and validation roles sit at the aggressive end of that trade by design. Second, validation layers are *multipliers*: the chamber runs several roles per analysis, so a per-role starvation rate compounds into per-analysis attempt inflation. The reactive cell’s debate IAF of 1.698 means the chamber’s effective price was $1.6\text{--}1.7\times$ its architecture-diagram price all night, invisibly, while it returned correct verdicts.

The outer-ring escalation sharpens this. When the inner retry also runs at the starved budget (pure retry), most inner retries fail (69 of 88 unrecovered), the orchestrator re-asks, and the layer’s job count itself inflates $1.6\times$ (193 versus 110–119). The cell’s debate layer thus paid three ways — starved first attempts, starved inner retries, and outer re-asks — to reach a scoreboard equivalent to every other cell’s. Teams adding adversarial validation, self-critique, or judge layers to production pipelines should assume by default that those layers, not the primary agents, are where attempt amplification will concentrate, and should give their instruments layer identity so the concentration is visible.

9 Non-Stationarity and Economic Forecasting

The companion paper reported a registered falsifier firing: a pure re-sample at the unchanged budget recovered far more starved jobs than attempt-exchangeability allowed, roughly 19 of 20 on 2026-08-09 [1]. Tonight, under the identical configuration, the same mechanism recovered 5 of 16 agent-layer and 19 of 88 debate-layer starved jobs. The same-budget rescue rate moved from roughly 95% to roughly 22–31% in two days on the same provider, model, corpus, and code. Meanwhile the budget-adding retry rescued 82 of 83 debate-layer starvations tonight — the mechanism that carries its own remedy stayed strong while the mechanism that relies on provider-side variance collapsed.

Two conclusions, one epistemic and one economic. Epistemically, attempt outcomes are not only non-exchangeable (the falsifier result) but *non-stationary*: a single night’s rescue rate, in either direction, must not be generalized. We publish tonight’s counts beside the companion night’s counts and decline to average them; the honest unit of report is the dated count with its denominator. The unresolved provider-side mechanism — the companion paper hypothesized cache-state or serving-path effects; our per-attempt cache hit/miss splits are recorded but not yet analyzed — remains open and is now better instrumented than before.

Economically, non-stationarity prices the difference between prevention and recovery as forecasting risk. What prevention stabilizes is the *attempt count at a configured budget*: one attempt

per job, with the starvation residual (measured tonight at 1 in 127) as the forecast error on that count. Realized dollars still move with output length, cache participation, and provider list prices — prevention removes the attempt-count term from the bill’s variance, not the whole bill. A recovery-based configuration’s cost model additionally contains the rescue rate as a load-bearing variable — and that variable just moved by a factor of three between two nights. Worse, its failure mode compounds: when rescue rates drop, inner retries fail, outer rings engage, and both attempt counts and job counts inflate (Section 8). An architecture that meets its scoreboard through retries has an attempt count that is a function of provider weather; one that meets it through headroom has an attempt count that is a function of its own configuration. At month scale, the first adds a random variable to the bill; the second leaves the bill’s variance where every pipeline already carries it — in output length and prices.

10 Implications for AI Cost Accounting

Two bases, kept separate. This study’s spend has two honest totals. The ledger-computed total is \$3.694: every attempt priced at undiscounted direct list rates, cache-hit input priced at the full input rate — an upper bound with per-attempt attribution. The account-billed total is \$3.07: the provider balance moved \$4.74 to \$1.67 across the run — the observed aggregate account-billed amount under the controlled run conditions, attributable to nothing smaller than the run. The \$0.62 residual between the bases is a measured pricing residual, consistent with the provider’s prompt-cache discount structure; where provider pricing terms are independently known it may be decomposed among realized cache discounts and other provider-side pricing adjustments, which the per-attempt hit/miss splits support analyzing. Both totals are correct on their own basis; mixing them in one figure would be an accounting error, and we nearly committed it in an early draft of our own ledger document. Cost reporting for LLM systems should name its basis the way financial statements name accrual versus cash.

Meter attempts, not responses. Any team whose pipeline contains a retry, a fallback, or a re-asking orchestrator already has a ledger-versus-legacy gap; ours was 9.6% of spend on a calm night in a tuned configuration, and the pure-retry cell shows the same pipeline carrying a 48.9% discard share in one layer under a plausible misconfiguration. The instrument is not expensive: one append per HTTP attempt at the client boundary, below every wrapper that can rebind a response.

Account for reasoning tokens as a first-class category. Reasoning tokens were 65.1% of this study’s billed output and 228,889 of its 228,890 discarded output tokens. A cost dashboard that reports visible completion tokens misses the majority of the bill; a scoreboard misses it by construction. The starved attempt is the extreme case: a bill consisting entirely of tokens no human or program will ever see.

Report iaf and eci alongside accuracy. Open-weight models and modern serving stacks have made inference increasingly *controllable*; the complementary requirement this paper argues for is that it be *accountable* — as pipeline value migrates into harnesses, procedures, verification, and provenance, the attempt ledger belongs to that layer. Both metrics are cheap to compute from an attempt ledger, both have measured denominators, and both separate configurations that accuracy cannot. A benchmark row of the form “47.0/48, 9/10 execution-clean, agent IAF 1.00, debate IAF 1.70, debate ECI 1.58” says what the first two entries hide: which machine produced the score, and what the score cost.

11 Limitations

The boundaries mirror the companion study’s, with economic additions. One production pipeline, one provider and model on the direct channel, one night per cell, $n=10$ per cell; the benchmark corpus and answer keys are the companion paper’s frozen but developer-written artifacts, so nothing here is a model-capability estimate. All cross-night comparisons in this paper are between two nights only, which establishes non-stationarity but not its distribution; rescue rates here bound nothing beyond the nights reported. ECI as computed is a mild upper bound (salvage paths are classified by reply emptiness, not downstream value; Section 3), and ledger dollars are upper bounds (cache-hit input at full rate). The account-billed total is inferred from balance deltas, not vendor token-level telemetry. The layer tag derives from the calling site’s naming convention, which is robust in this pipeline but is a convention, not a protocol. The outer-ring job-inflation signature (193 versus 110–119) is measured, and its attribution to orchestrator re-asks follows from the client’s identity design, but we did not add a separate orchestrator-side instrument to cross-check it; a second ring-level instrument is the obvious next hardening. The per-attempt cache hit/miss splits bearing on the unresolved rescue-rate mechanism are recorded but not analyzed here. Finally, list prices and cache-discount policies are provider facts current as of the run dates and will drift.

12 Conclusion

The companion study argued that a benchmark score from a fallback-enabled pipeline is un-attributable without degradation accounting. This paper makes the parallel argument for the invoice: a cost figure from a retry-enabled pipeline is unattributable without attempt accounting. One ledger row per physical attempt — with logical-job identity, layer identity, reasoning-token counts, and outcome — was sufficient to measure what the previous study could only count: prevention measured as a counted starvation residual with no hidden recovery amplification, rather than inferred from a silent log; a validation layer quietly running at 1.7 times its diagrammed price; half a layer’s spend buying discarded thinking under a plausible misconfiguration; and a rescue rate that moved by a factor of three in two days under frozen configuration. All of it sat beneath four scoreboard-equivalent cells.

The scoreboard records what the pipeline answered. The ledger records what the pipeline did. A retry repairs the answer. It does not erase the first invoice.

Data and Artifact Availability

The run is fully committed in the project repository: the run-of-record script (four cells, lever hygiene, balance-floor guard), the four raw per-attempt JSONL ledgers, the twenty per-launch result files carrying per-repetition boards and agent modes, and the compiled report (`econ_report_2026-08-10.js`) from which every number in this paper is drawn. The provenance chain for every figure is short and mechanical: raw JSONL ledgers $\rightarrow$ per-launch result JSONs $\rightarrow$ analyzer $\rightarrow$ compiled report $\rightarrow$ manuscript tables and figure, with the figure script reading the report programmatically rather than transcribing values. The analyzer itself is a direct implementation of the definitions given in this paper and is retained privately as part of the measurement implementation (see the patent notice in the repository); its SHA-256 digest is recorded in the repository’s published manifest, and every reported quantity is independently recomputable from the published ledgers

and result files using those definitions. The run-of-record script is provenance, not a portable reproducer: it encodes the machine-local repository path and environment of the original run so that the artifact records exactly what executed; re-running the study elsewhere requires the project repository and provider credentials. The instrumented client and the accompanying ledger document are in the same history (instrumentation commit and a correction commit fixing two aggregate-dollar errors caught in review — an accounting discrepancy found while writing a paper about accounting, which we record in the spirit of the method). The frozen artifact bundle whose SHA-256 digest is recorded in the related provisional patent application is retained privately as evidence; the public repository carries the data, the manifest, and this manuscript.

References

- [1] S. Halverson. When the Model Isn’t the Problem: Degradation Accounting in Fallback-Enabled Production LLM Pipelines. Preprint v0.13.1, VuduVations, August 2026.
- [2] OpenAI. Reasoning models — API documentation (reasoning tokens are billed as output tokens from the completion budget). <https://platform.openai.com/docs/guides/reasoning>, accessed 2026-08.
- [3] DeepSeek. Models & Pricing; Thinking Mode — API documentation. <https://api-docs.deepseek.com>, accessed 2026-08.
- [4] L. Chen, M. Zaharia, and J. Zou. FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance. arXiv:2305.05176, 2023.
- [5] I. Ong et al. RouteLLM: Learning to Route LLMs with Preference Data. arXiv:2406.18665, 2024.
- [6] C. Snell, J. Lee, K. Xu, and A. Kumar. Scaling LLM Test-Time Compute Optimally Can Be More Effective than Scaling Model Parameters. arXiv:2408.03314, 2024.
- [7] J. Dean and L. A. Barroso. The Tail at Scale. *Communications of the ACM*, 56(2):74–80, 2013.
- [8] B. Beyer, C. Jones, J. Petoff, and N. R. Murphy, editors. *Site Reliability Engineering: How Google Runs Production Systems*. O’Reilly Media, 2016.
- [9] D. Sculley et al. Hidden Technical Debt in Machine Learning Systems. In *NeurIPS*, 2015.
- [10] P. Liang et al. Holistic Evaluation of Language Models. arXiv:2211.09110, 2022.